

# RNNs Notes Final

Sunday, October 17, 2021 7:54 AM

<https://www.kaggle.com/kcsener/8-recurrent-neural-network-rnn-tutorial>

<https://www.kaggle.com/prashant111/comprehensive-guide-to-rnn-with-keras>

<https://www.kaggle.com/shivamb/beginners-guide-to-text-generation-using-lstms>

<https://towardsdatascience.com/beautifully-illustrated-nlp-models-from-rnn-to-transformer-8>

**n-gram Language Model**

**Suppose we are learning to predict the next word in a sentence.**

[0d69faf2109](#)

## Language Models: Exam

ining a 4-gram Language Model

ople

odel.

~~as the proctor stands~~

discards

$P(w|\text{students open})$

For example, suppose that

- “students opened the door”
- “students opened the window”

~~started the clock, the student~~  
rd

conc

$$\text{opened their}) = \frac{\text{count(stud)}}{\text{count(stu}}}$$

that in the corpus:

“eir” occurred 1000 times

“eir books” occurred 400 tim

nts opened their \_\_\_\_\_

dition on this

lents opened their *w*)  
udents opened their)

es

- $\rightarrow P(\text{books} \mid \text{student})$
- “students opened their books”
  - $\rightarrow P(\text{exams} \mid \text{student})$

9

## Sparsity Problem

### Sparsity Problem 1

**Problem:** What if “student opened their  $w$ ” never occurred in data? Then

lents opened their) = 0.4

their exams" occurred 100 times

lents opened their) = 0.1

## ns with n-gram Language Models

students  
r  
n w



(Partial) Solution  
to the count problem  
This is called

nes

## Language Models

Solution: Add small  $\delta$   
for every  $w \in V$ .

has probability 0!

$P(\mathbf{w}|\text{students open})$

## Sparsity Problem 2

**Problem:** What if “students opened their” never occurs in the data? Then we can’t calculate the probability for *any w*!

This is called

ned their) =

count(students c

count(students

nts  
urred in  
culate

**(Partial) Solution**  
on “opened to”  
This is called

*smoothing.*

opened their *w*)

(s opened their)

tion: Just condition

“heir” instead.

*backoff.*

# A RNN Language

output distribution

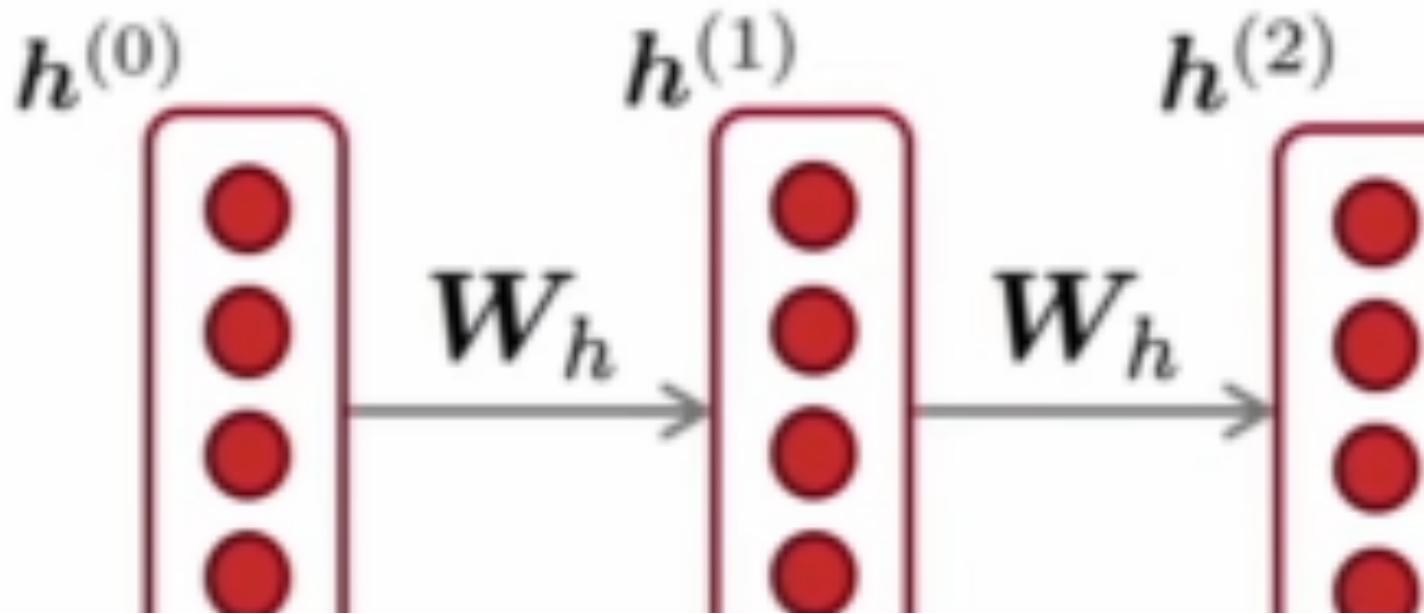
$$\hat{y}^{(t)} = \text{softmax} \left( U h^{(t)} + b_2 \right) \in \mathbb{R}^{|V|}$$

hidden states

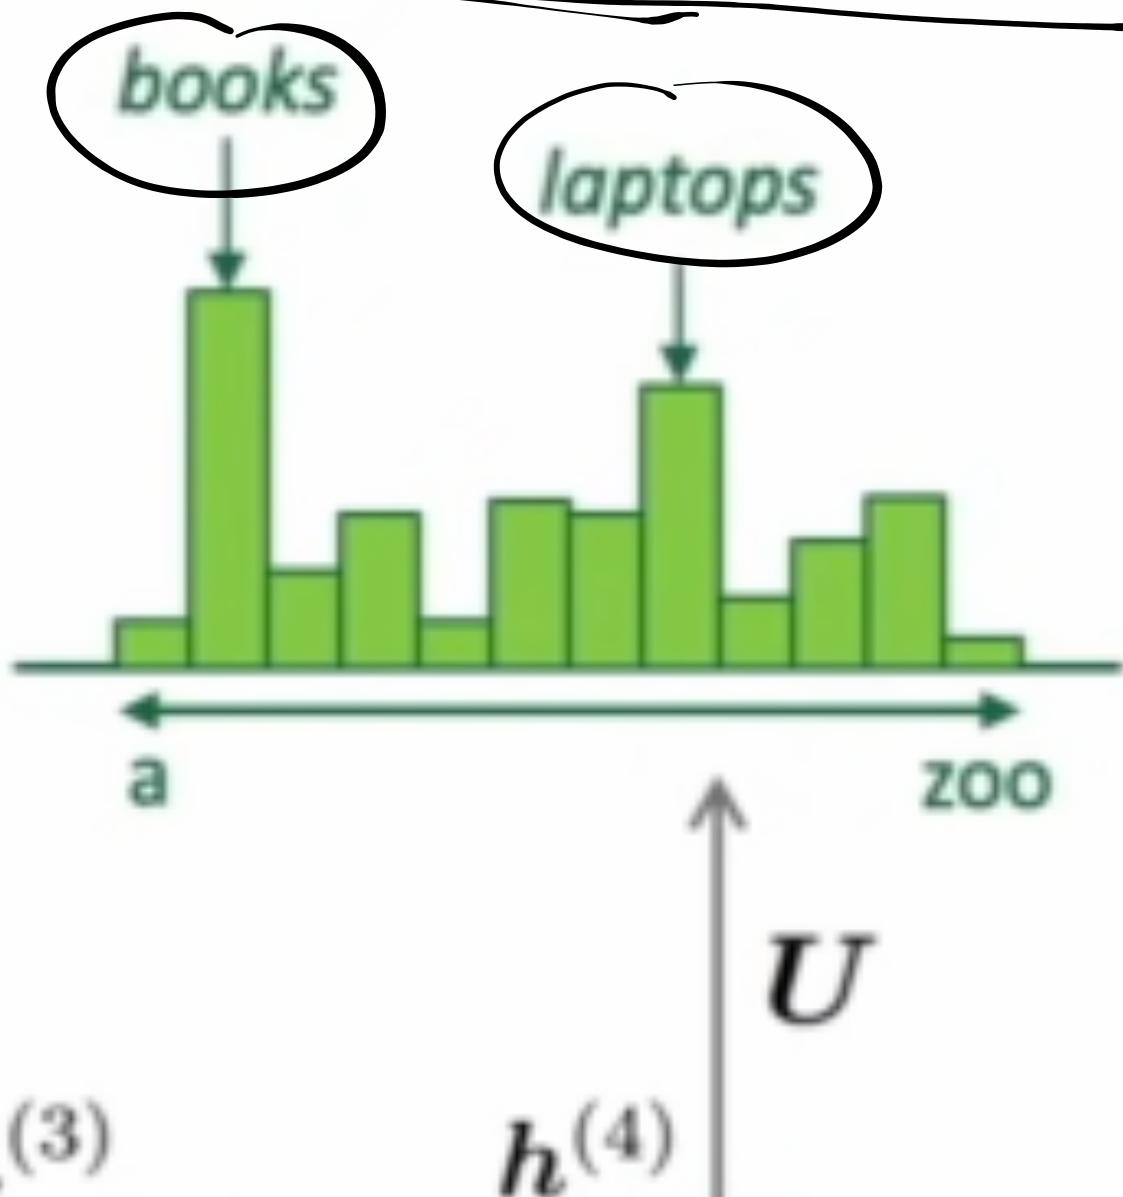
$$h^{(t)} = \sigma \left( W_h h^{(t-1)} + W_e e^{(t)} + b \right)$$

$h^{(0)}$  is the initial hidden state

# Image Model



$\hat{y}^{(4)} = P(\mathbf{x}^{(5)} | \text{the students opened their})$



$h^{(3)}$

$h^{(4)}$

$W_h$

$W_h$





## word embeddings

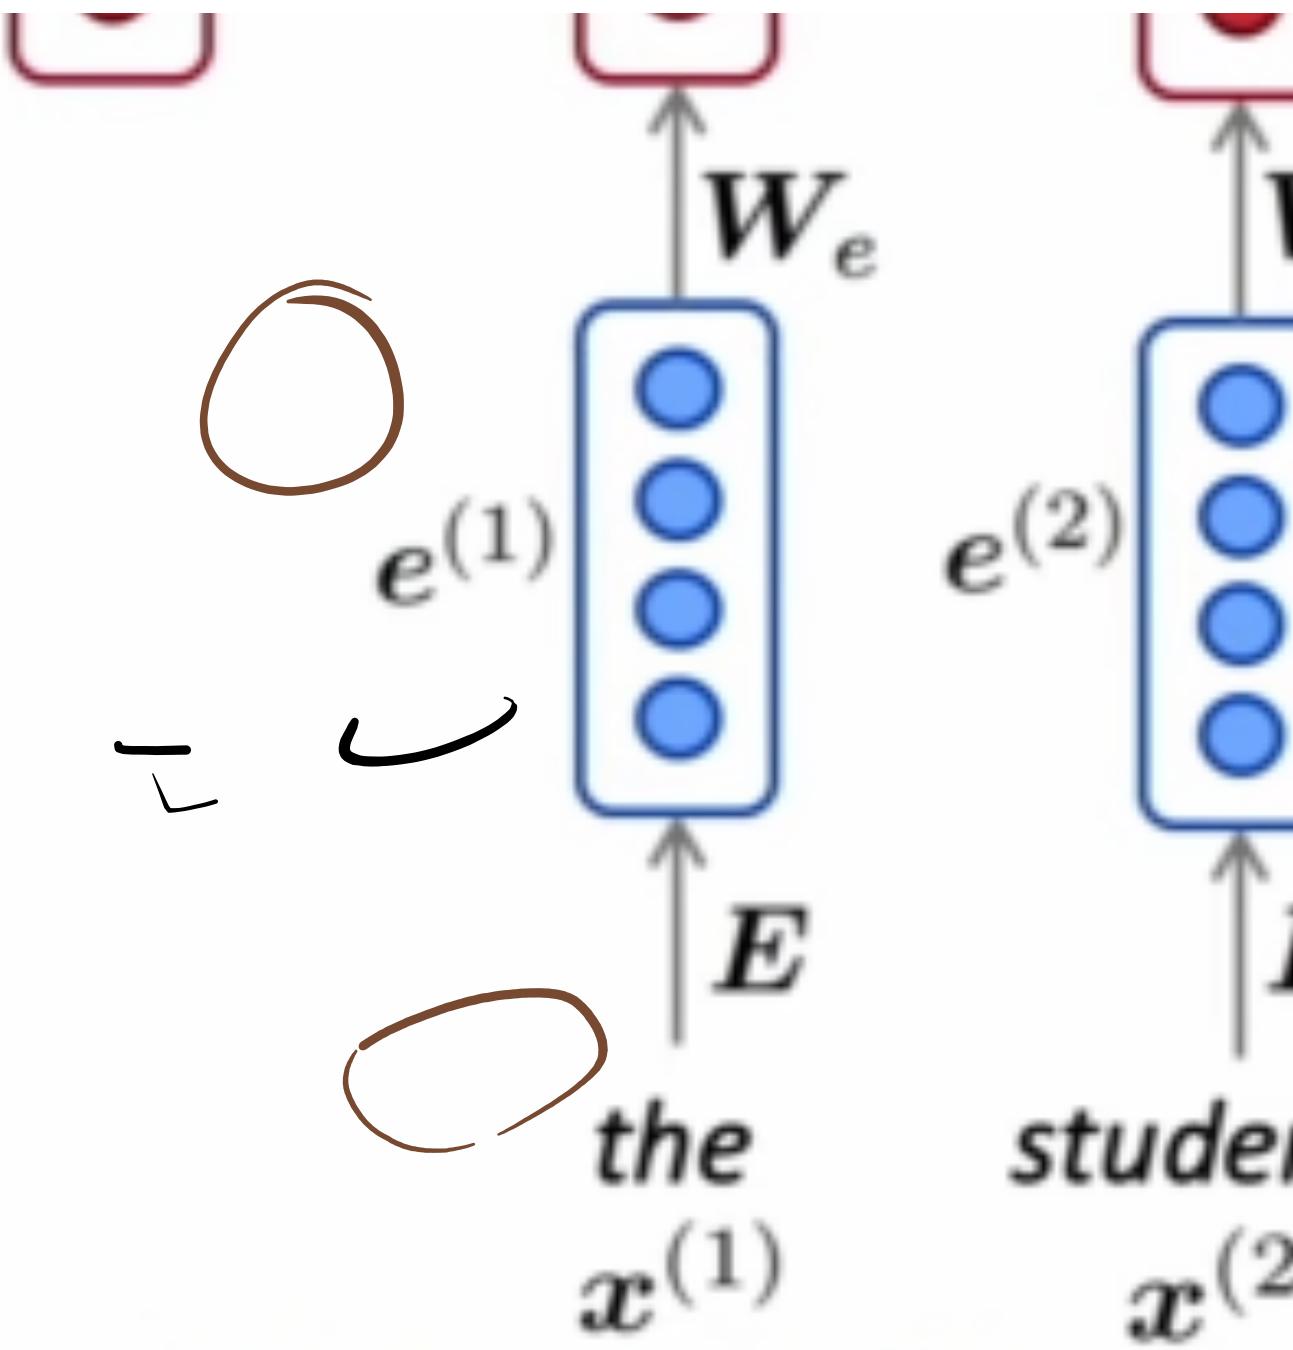
$$e^{(t)} = \mathbf{E}x^{(t)}$$

## words / one-hot vectors

$$x^{(t)} \in \mathbb{R}^{|V|}$$

23

Note:  
longer



this input sequence could be much  
longer, but this slide doesn't have space!

$W_e$

$E$

nts

(3)



$e^{(3)}$



*opened*

$x^{(3)}$

$e^{(4)}$



*their*

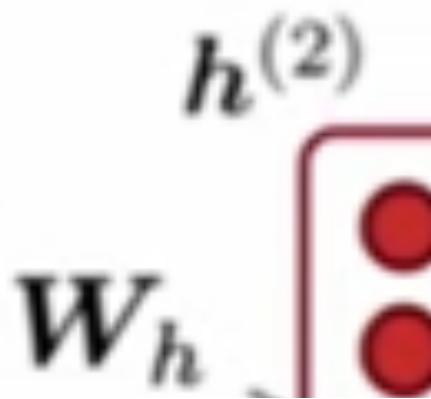
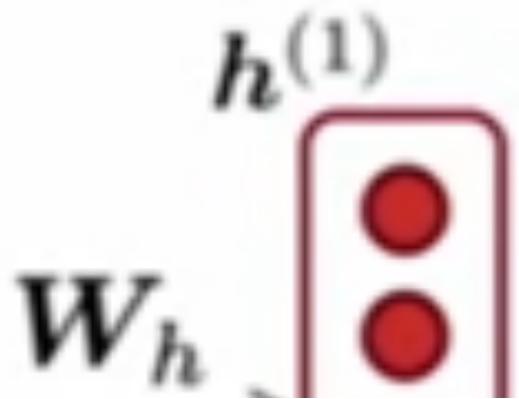
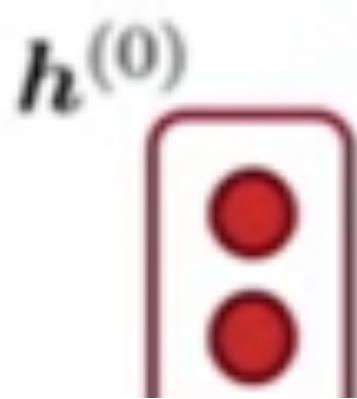
$x^{(4)}$

# A RNN Language

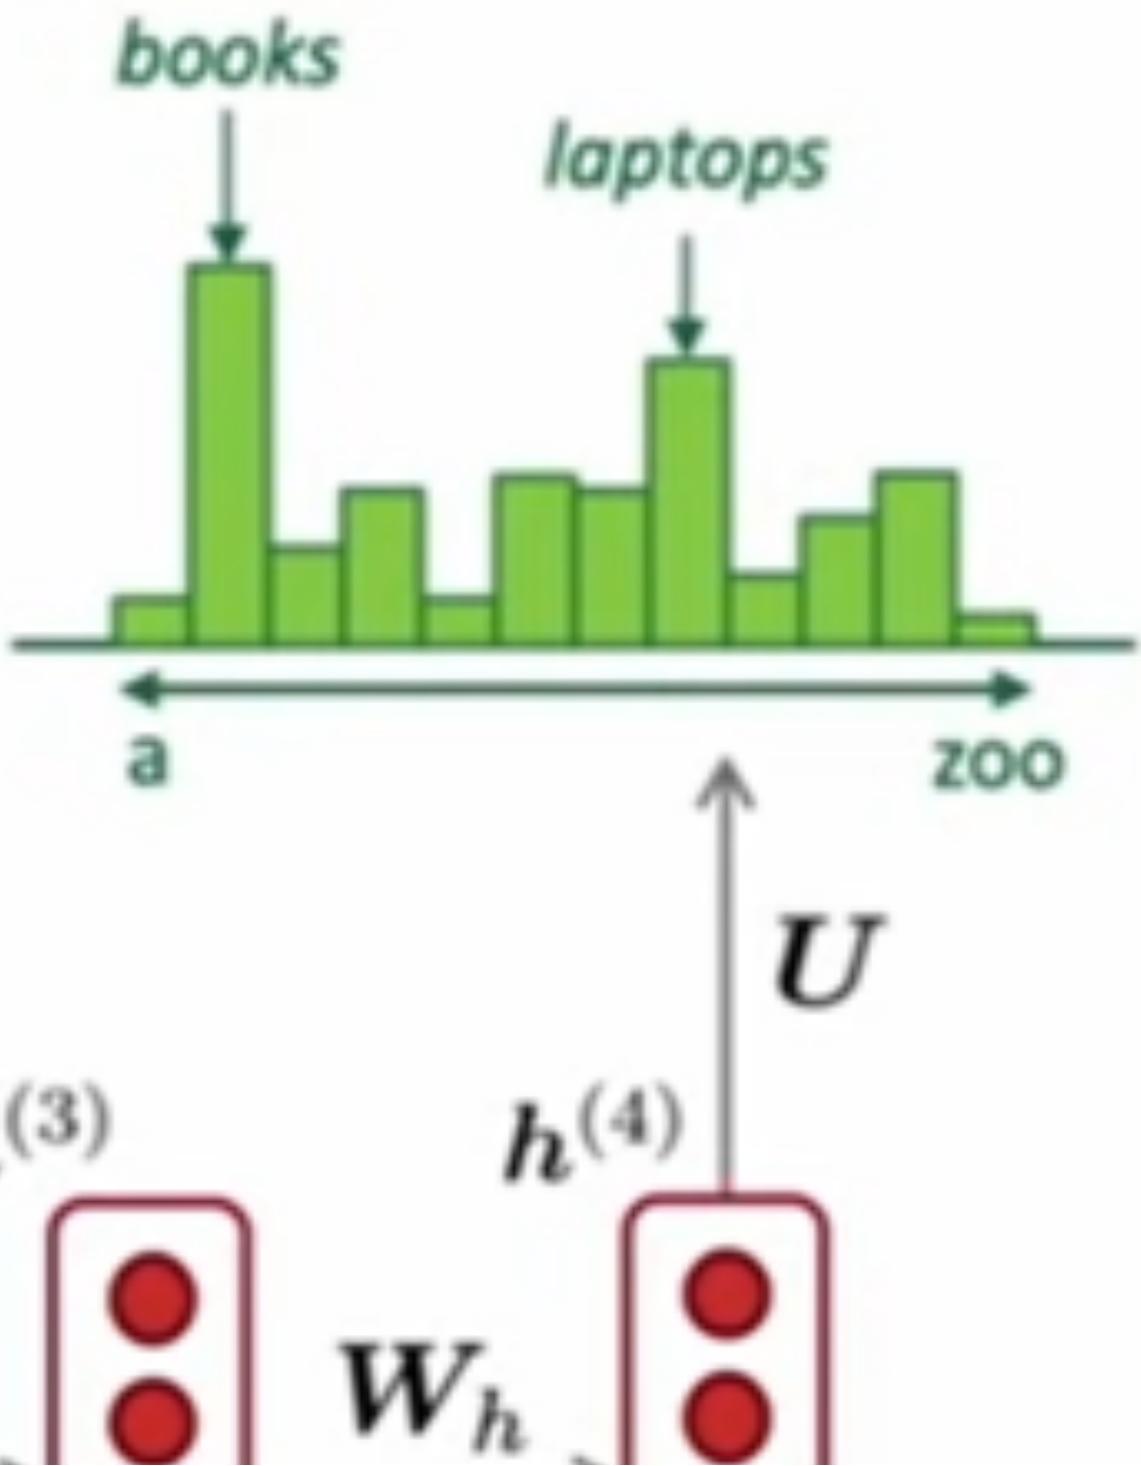
## RNN Advantages:

- Can process any length input
- Computation for step  $t$  can (in theory) use information from many steps back
- Model size doesn't

# Image Model



$\hat{y}^{(4)} = P(\mathbf{x}^{(5)} | \text{the students opened their}$



**increase** for longer input

- Same weights applied on every timestep, so there **symmetry** in how inputs are processed.

## RNN **Disadvantages:**

- Recurrent computation is **slow** 
- In practice, difficult to access information from **many steps back**

is

s  
meeney

—

$e^{(1)}$

$E$

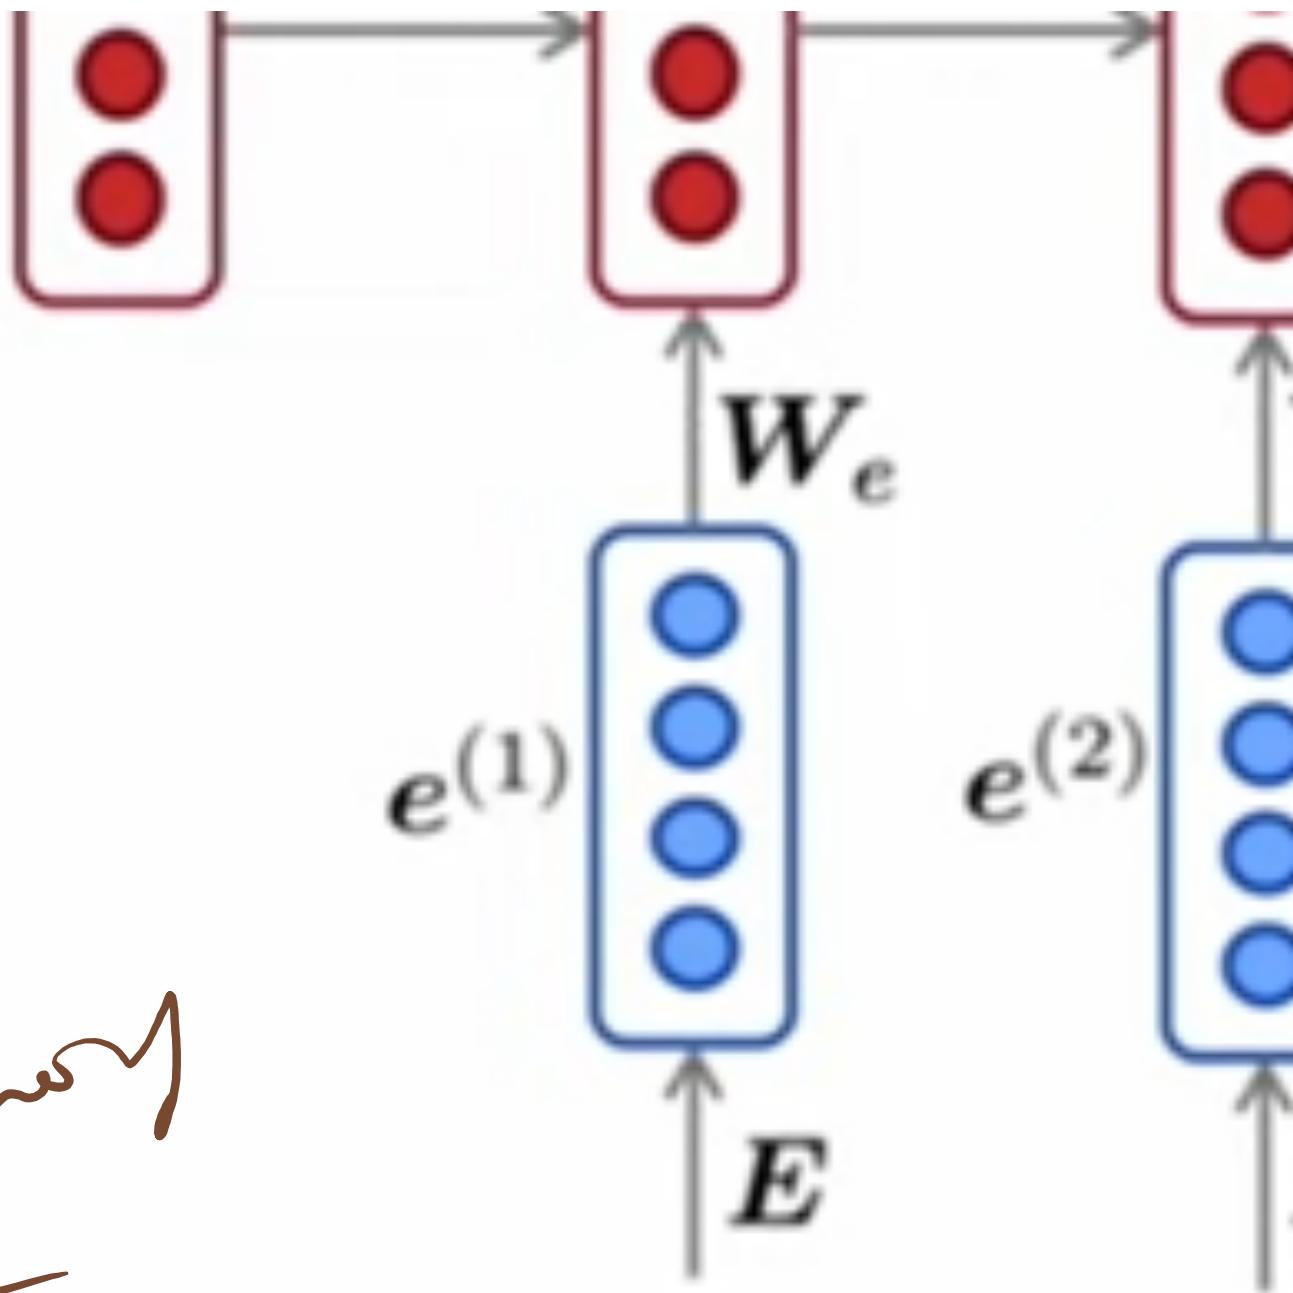
the

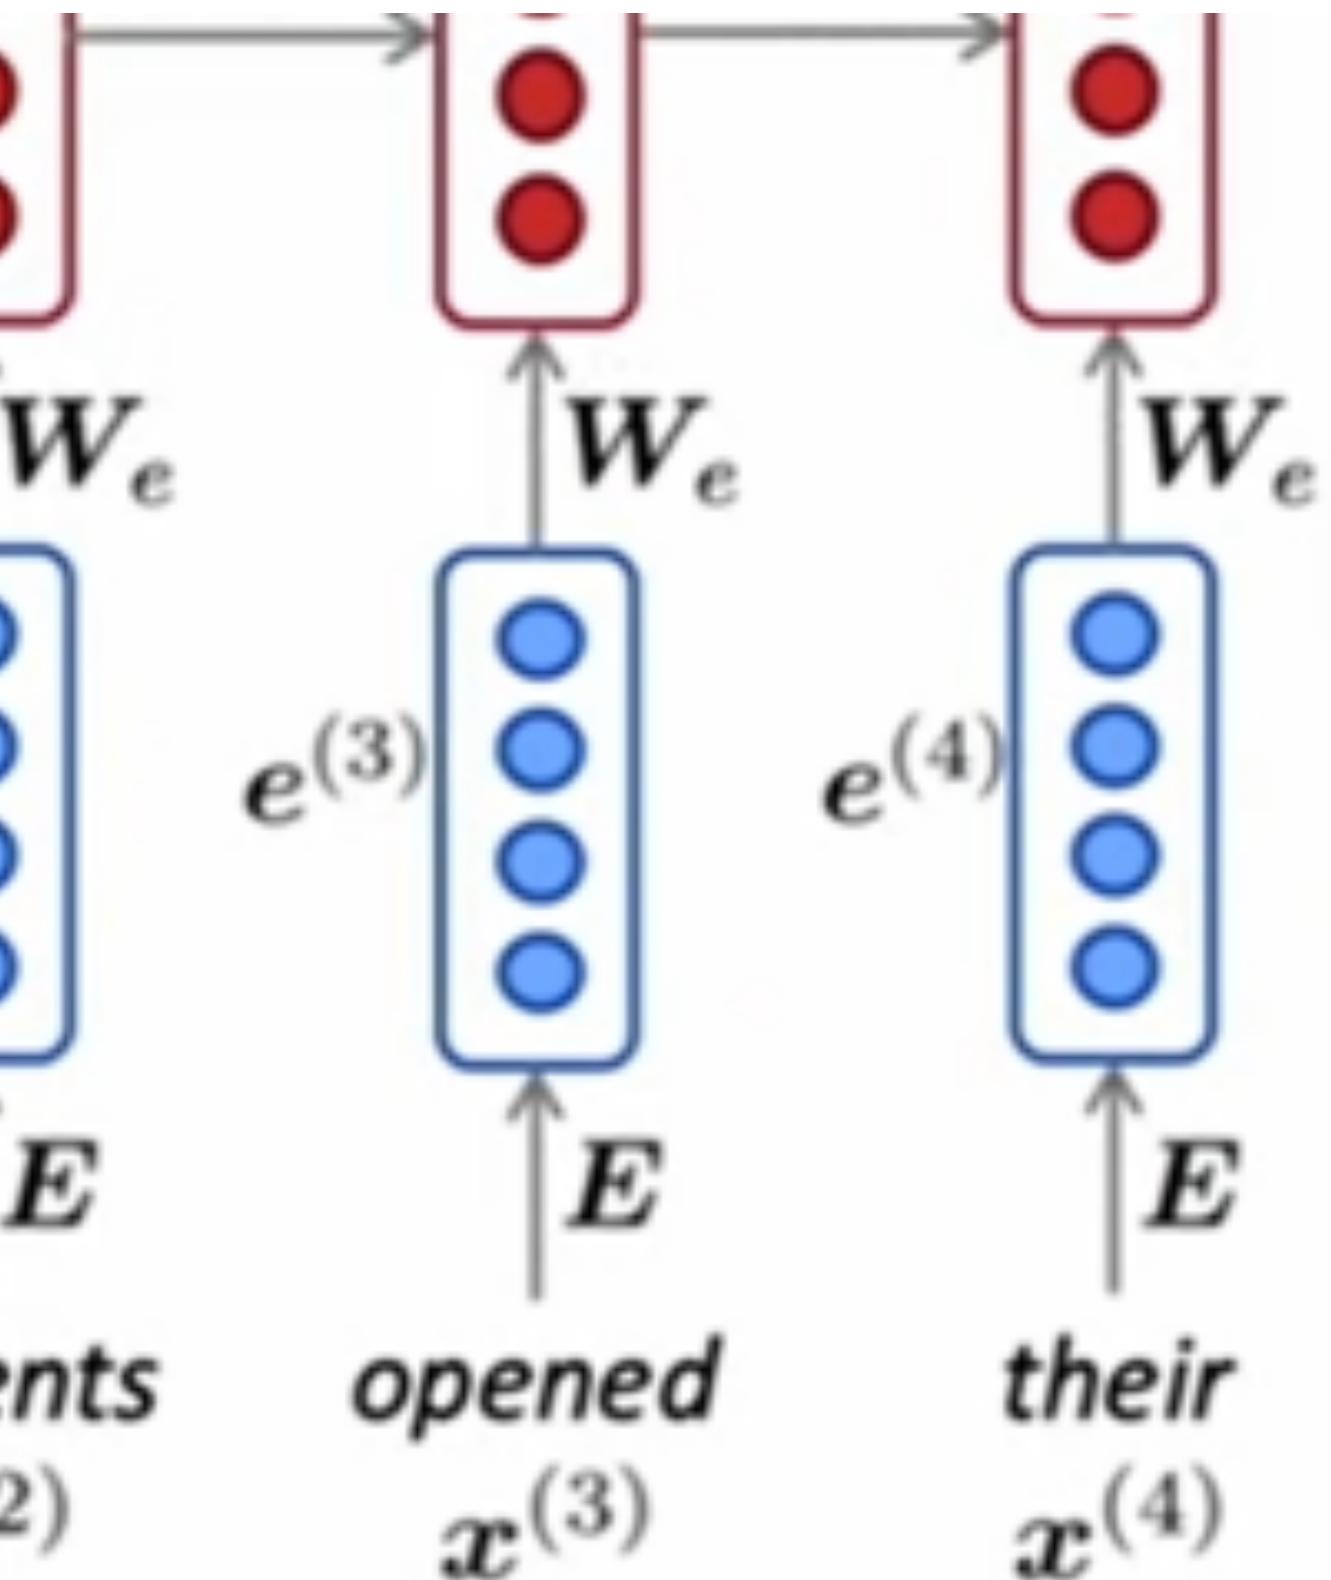
$x^{(1)}$

$e^{(2)}$

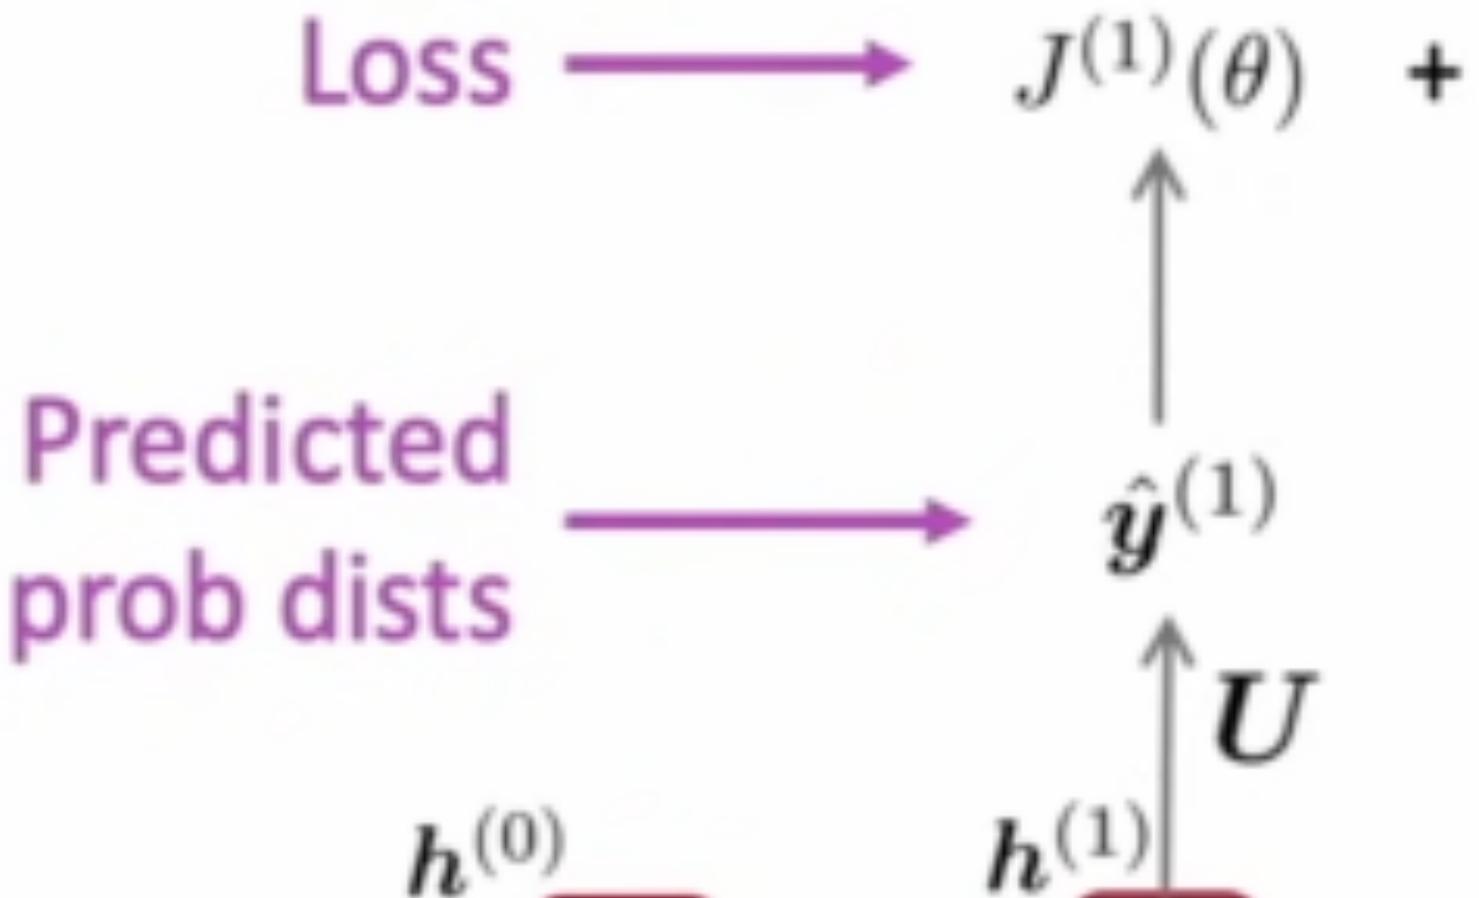
student

$x^{(2)}$





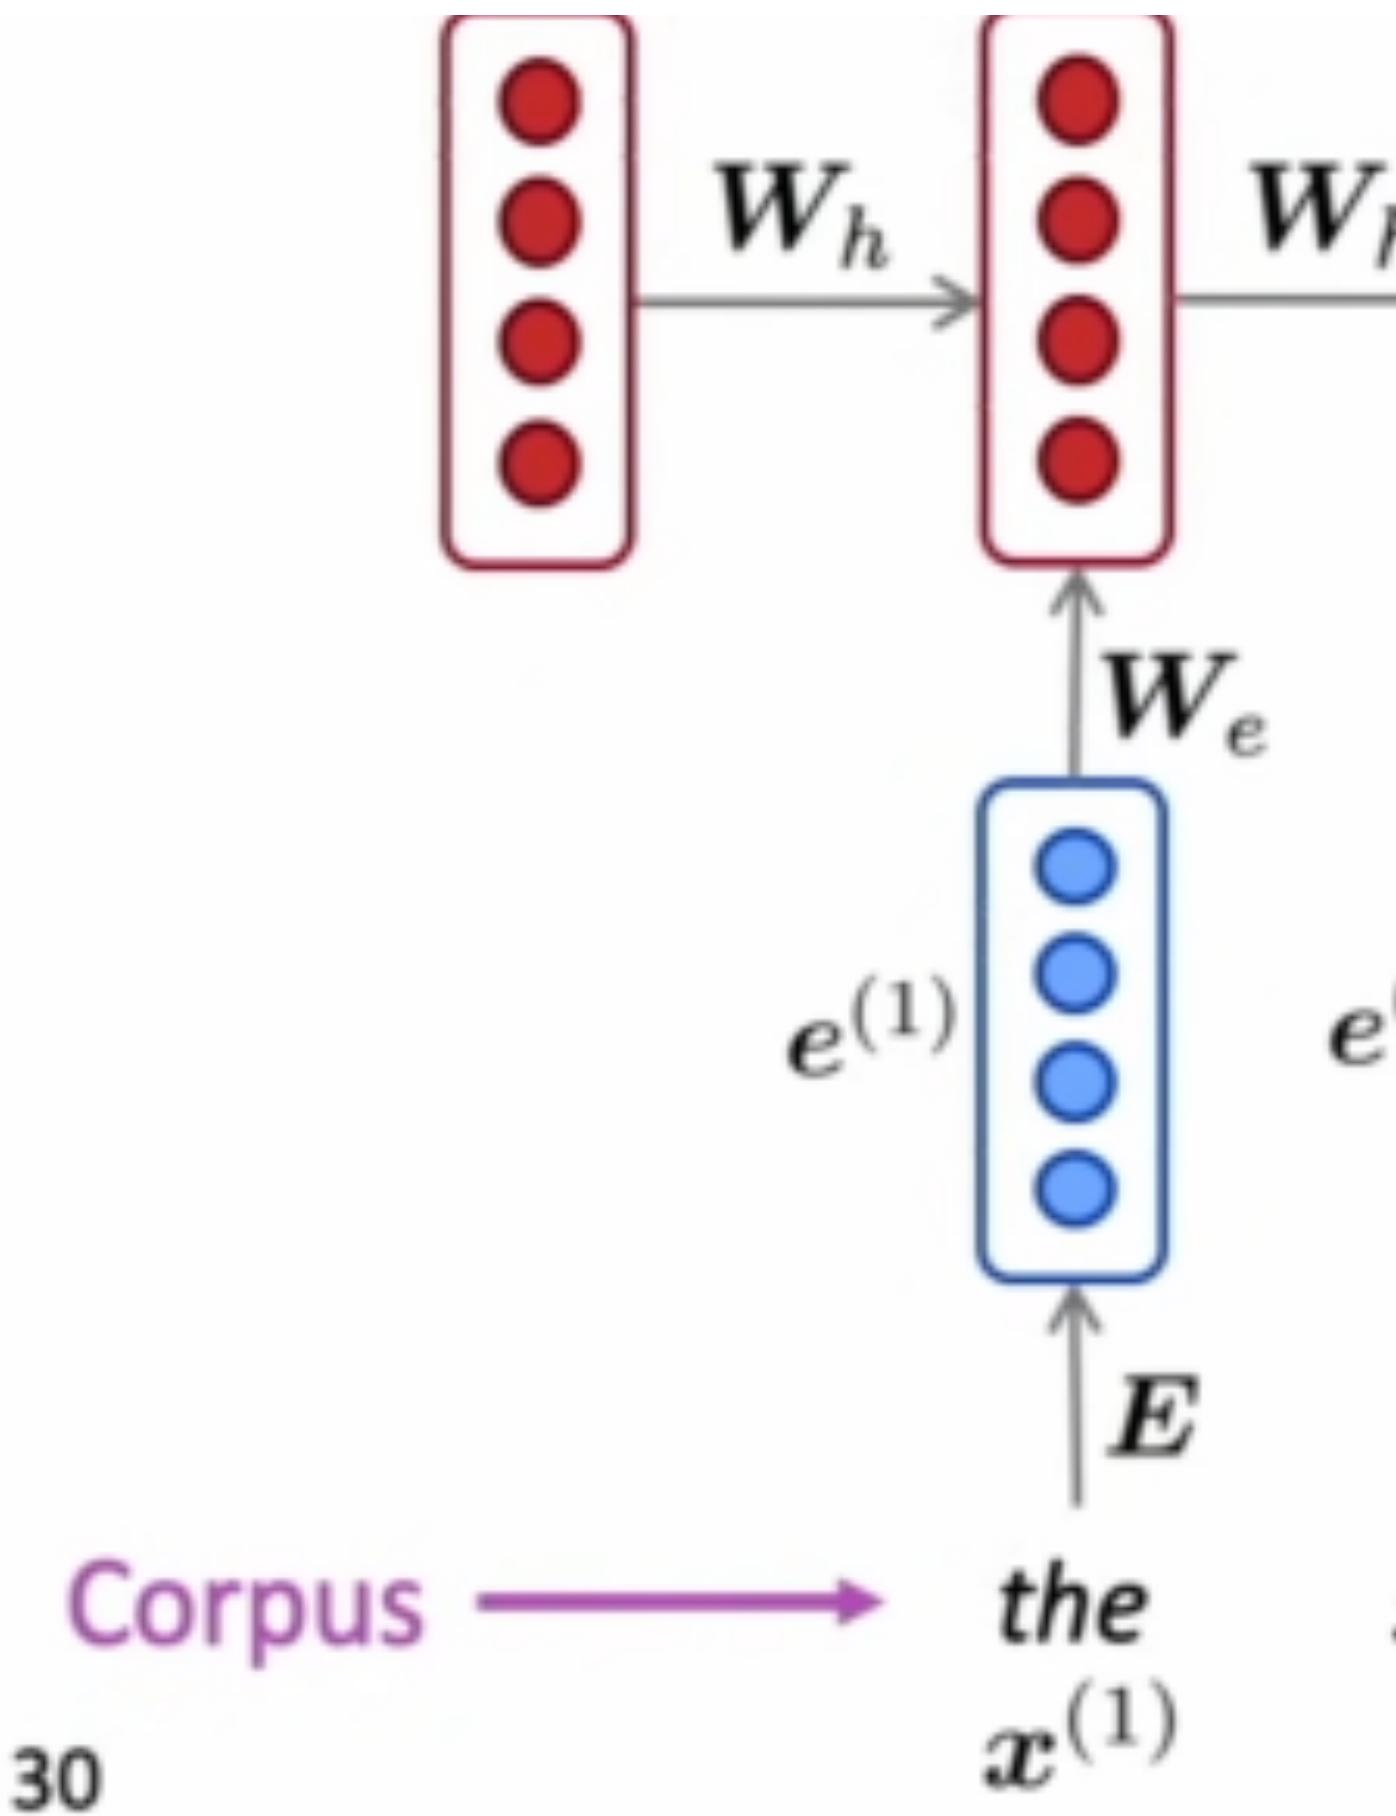
# Training a RNN L

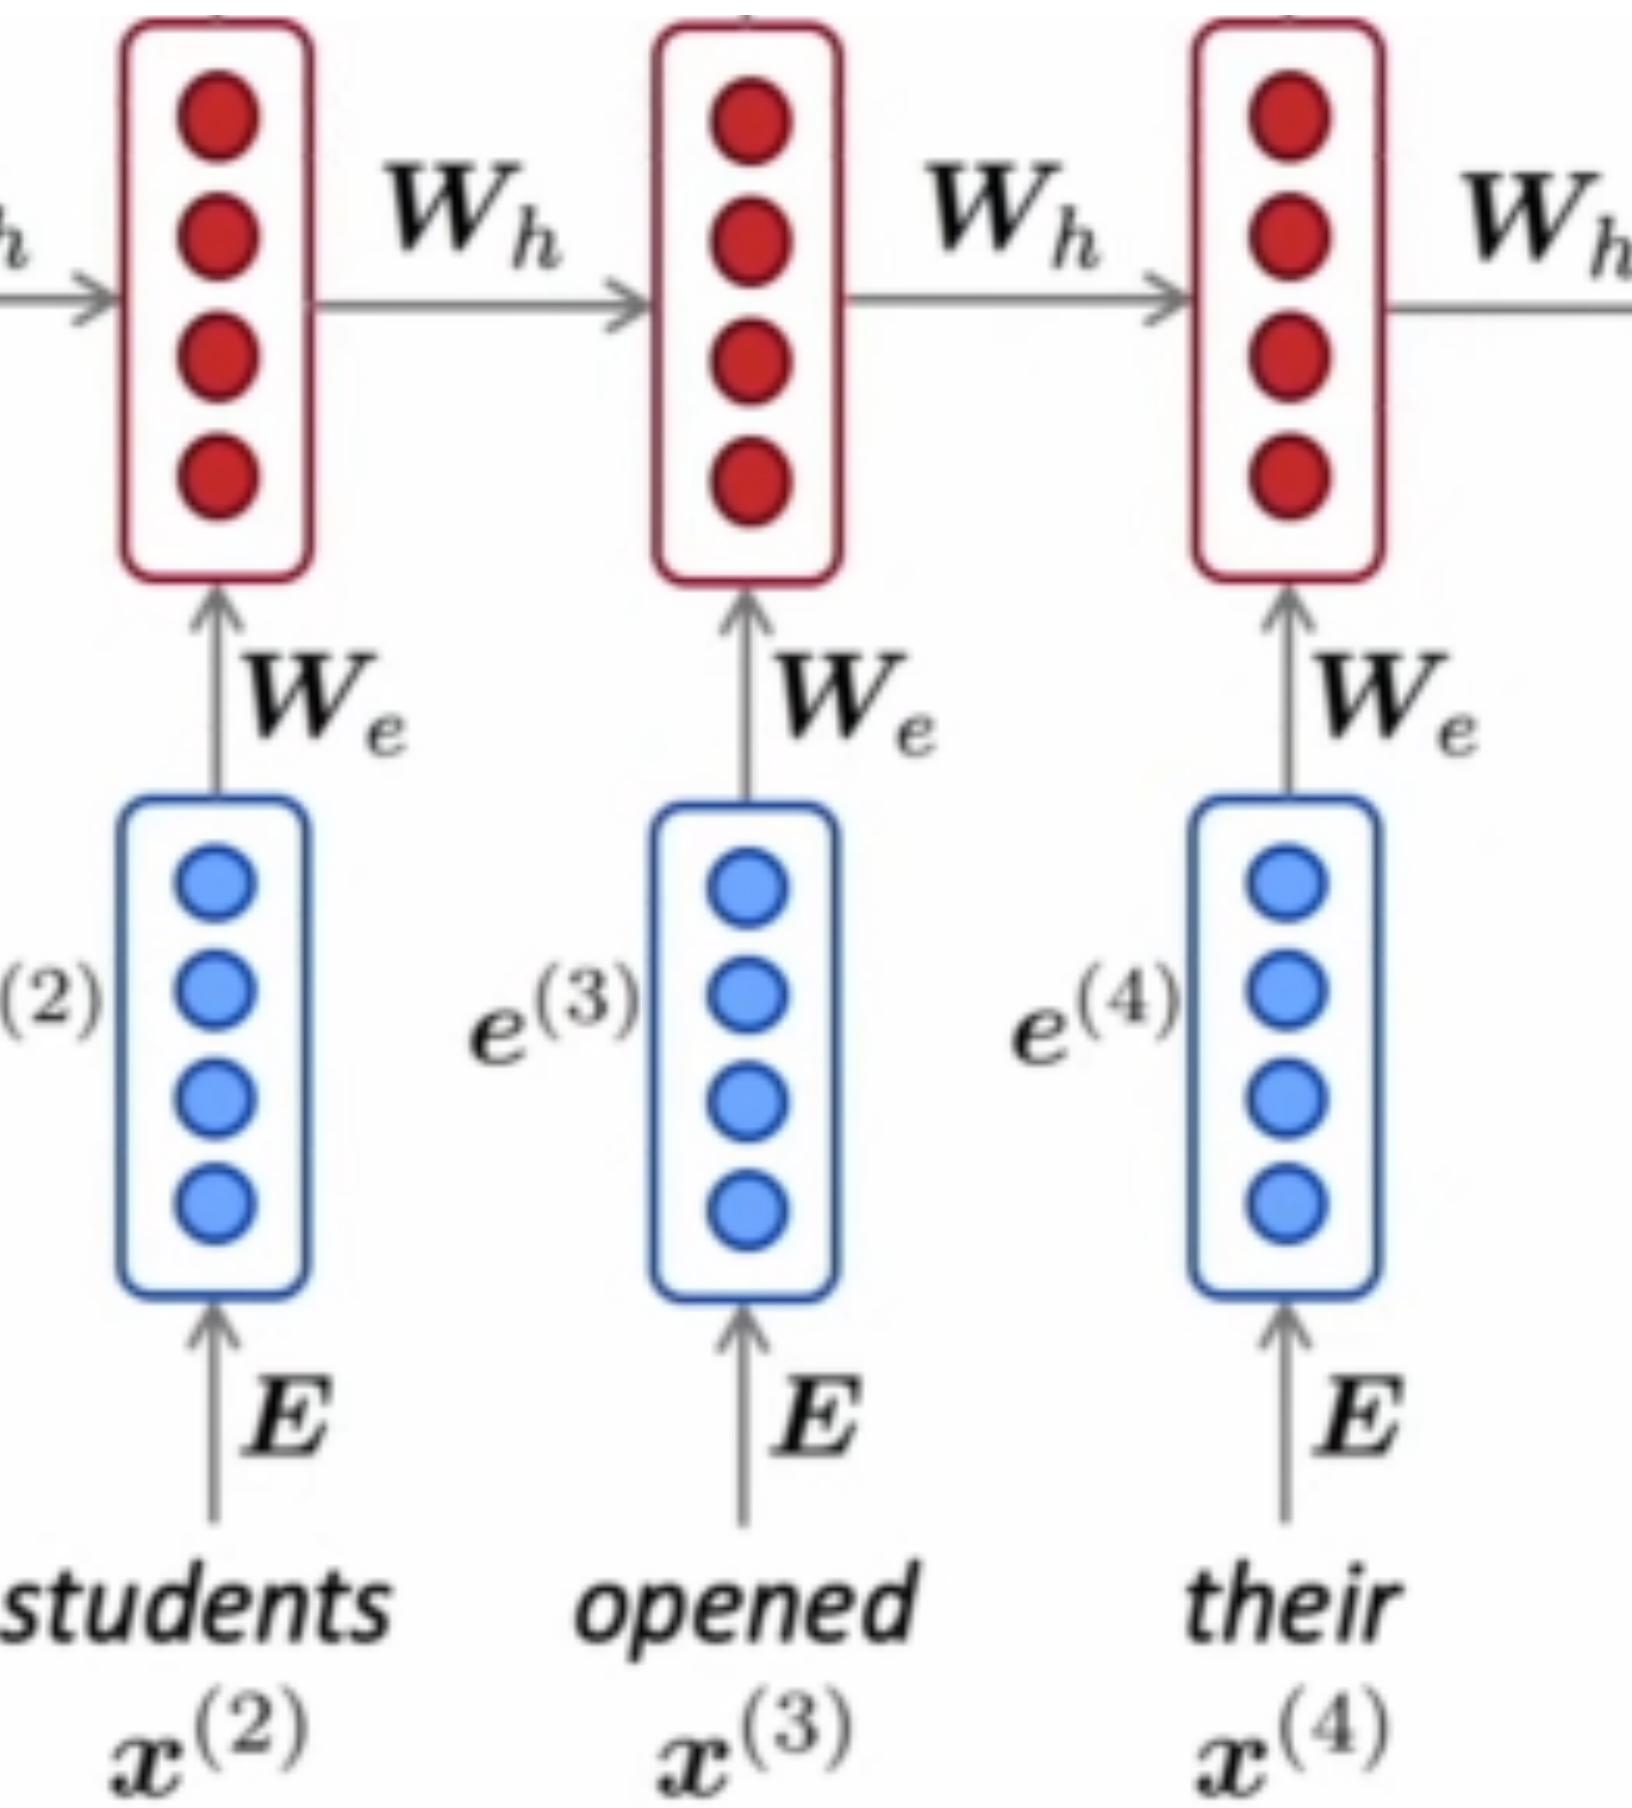


# Language Model

$$J^{(2)}(\theta) + J^{(3)}(\theta) + J^{(4)}(\theta) + \dots$$
$$\hat{y}^{(2)}$$
$$\hat{y}^{(3)}$$
$$\hat{y}^{(4)}$$
$$h^{(2)}$$
$$h^{(3)}$$
$$h^{(4)}$$
$$U$$
$$U$$
$$U$$

$$... = J(\theta) = \frac{1}{T} \sum_{t=1}^T J^{(t)}(\theta)$$







...  
...

*exams*

...

# Generating text

Just like a n-gram Language Model, we can generate text by repeatedly predicting the next word.

*favorite*



sample

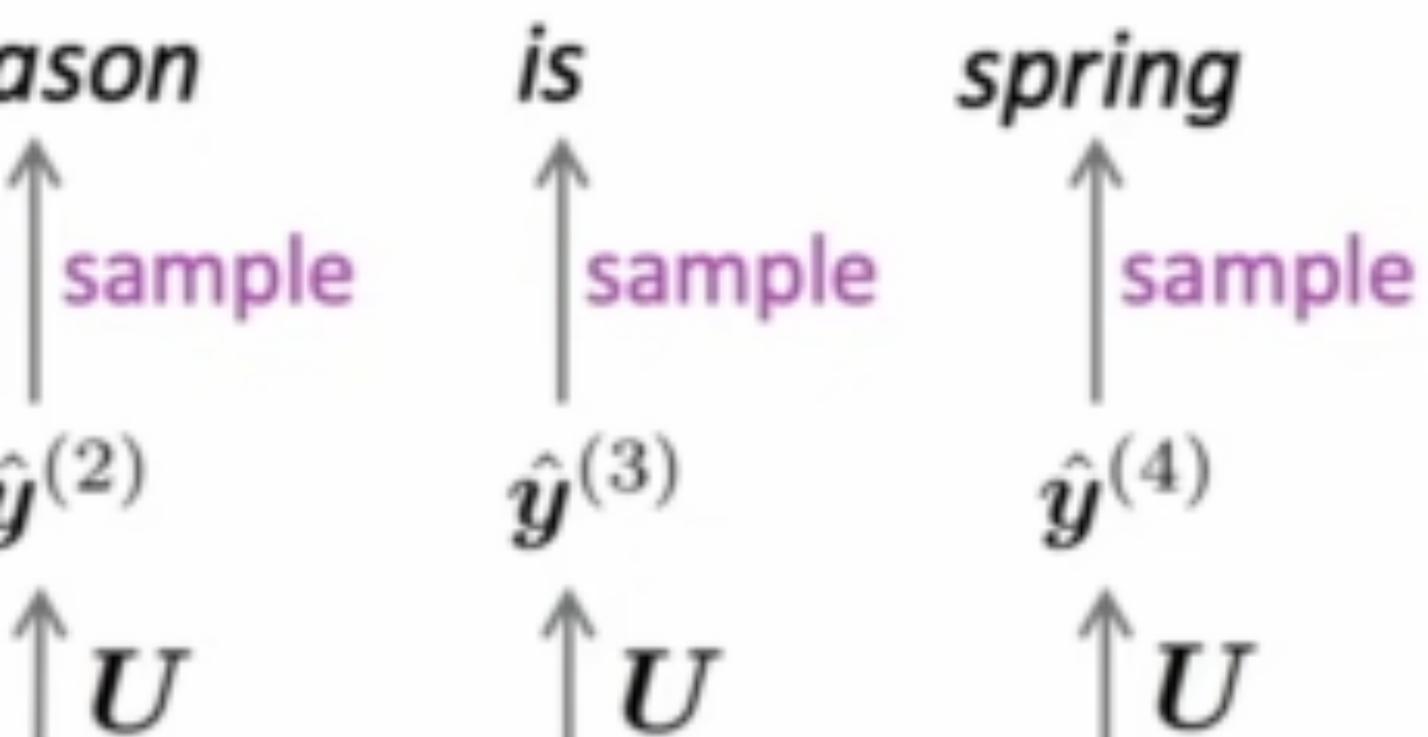
$\hat{y}^{(1)}$



$U$

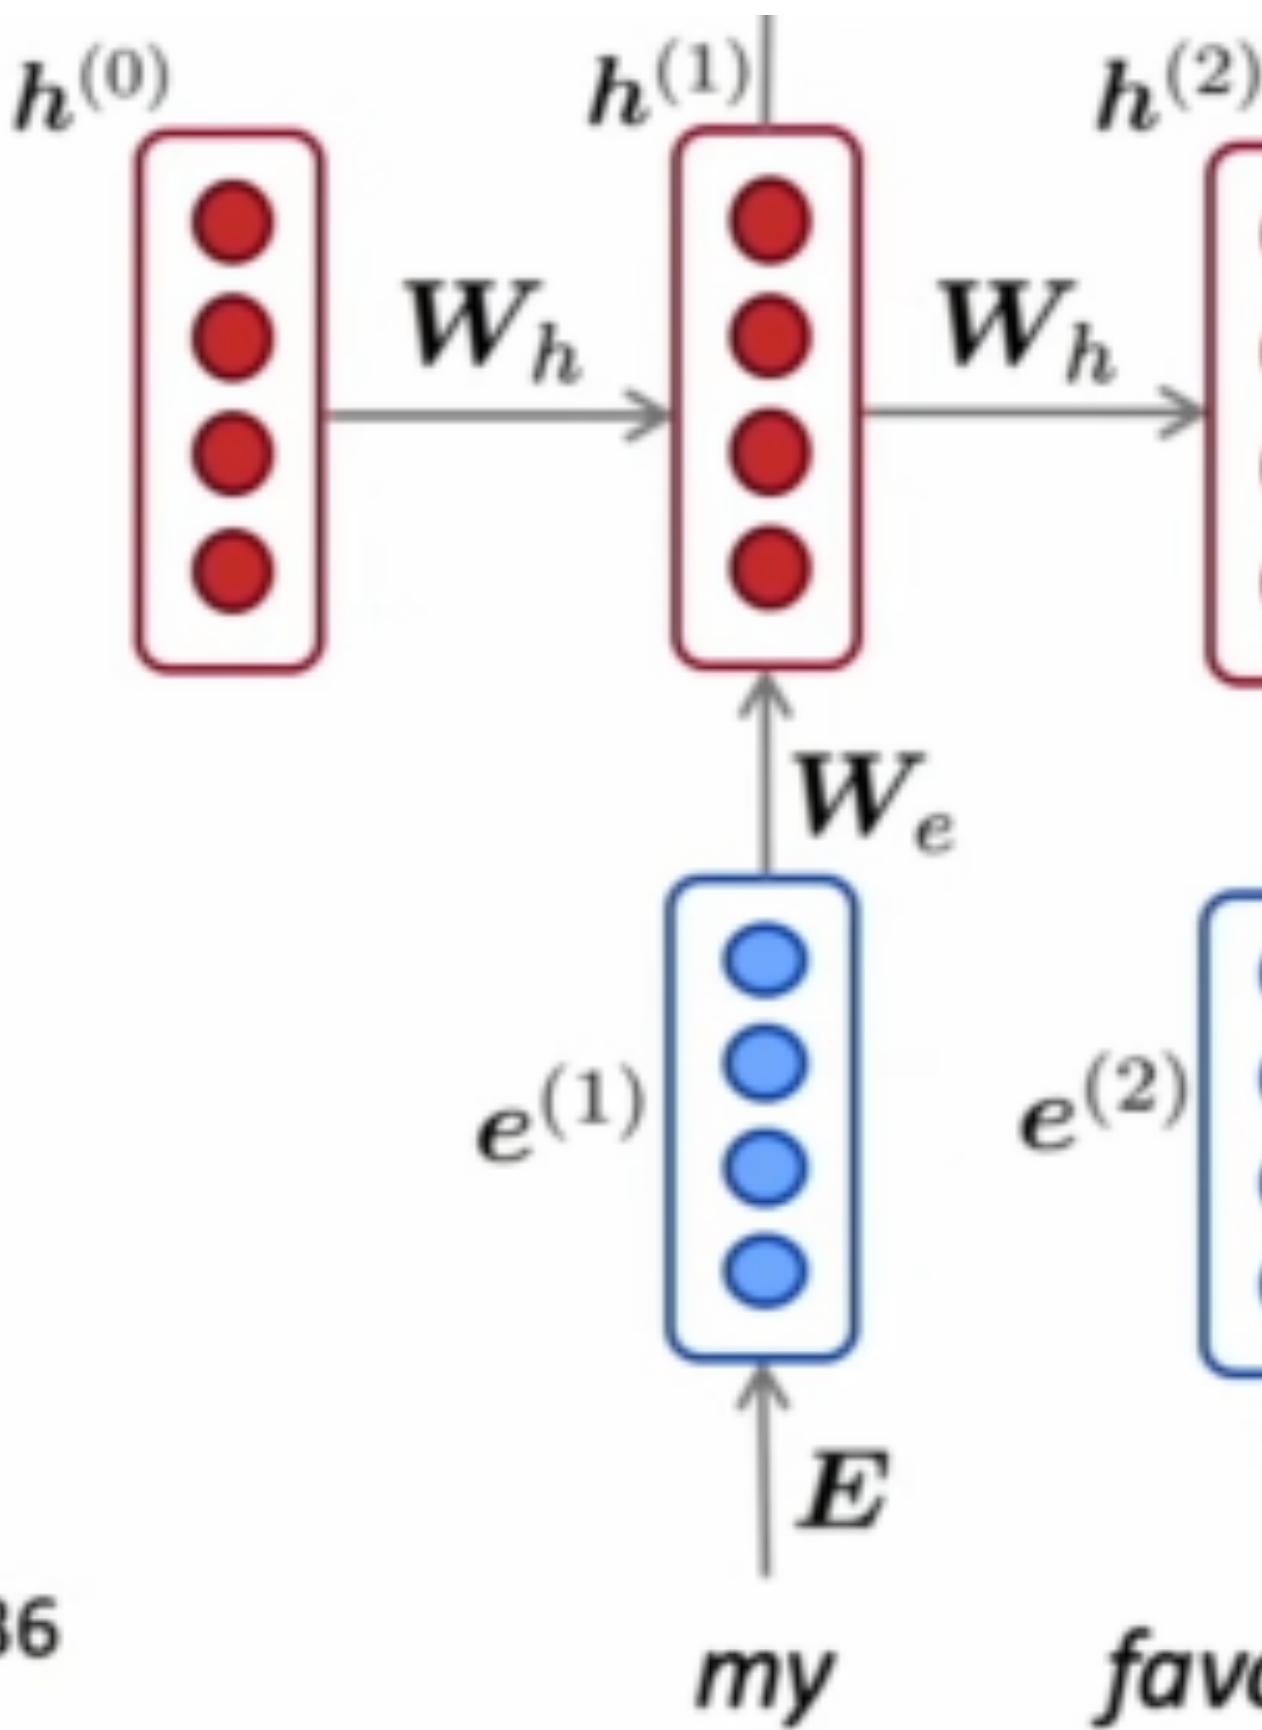
# Start with a RNN Language Model

In this section, we will learn how to build a Language Model, you can use a RNN Language Model to predict the next word in a sentence. We will use a simple RNN Language Model and sampling. Sampled output is not necessarily the most likely word, but it is more interesting than a deterministic model.



# Language Model

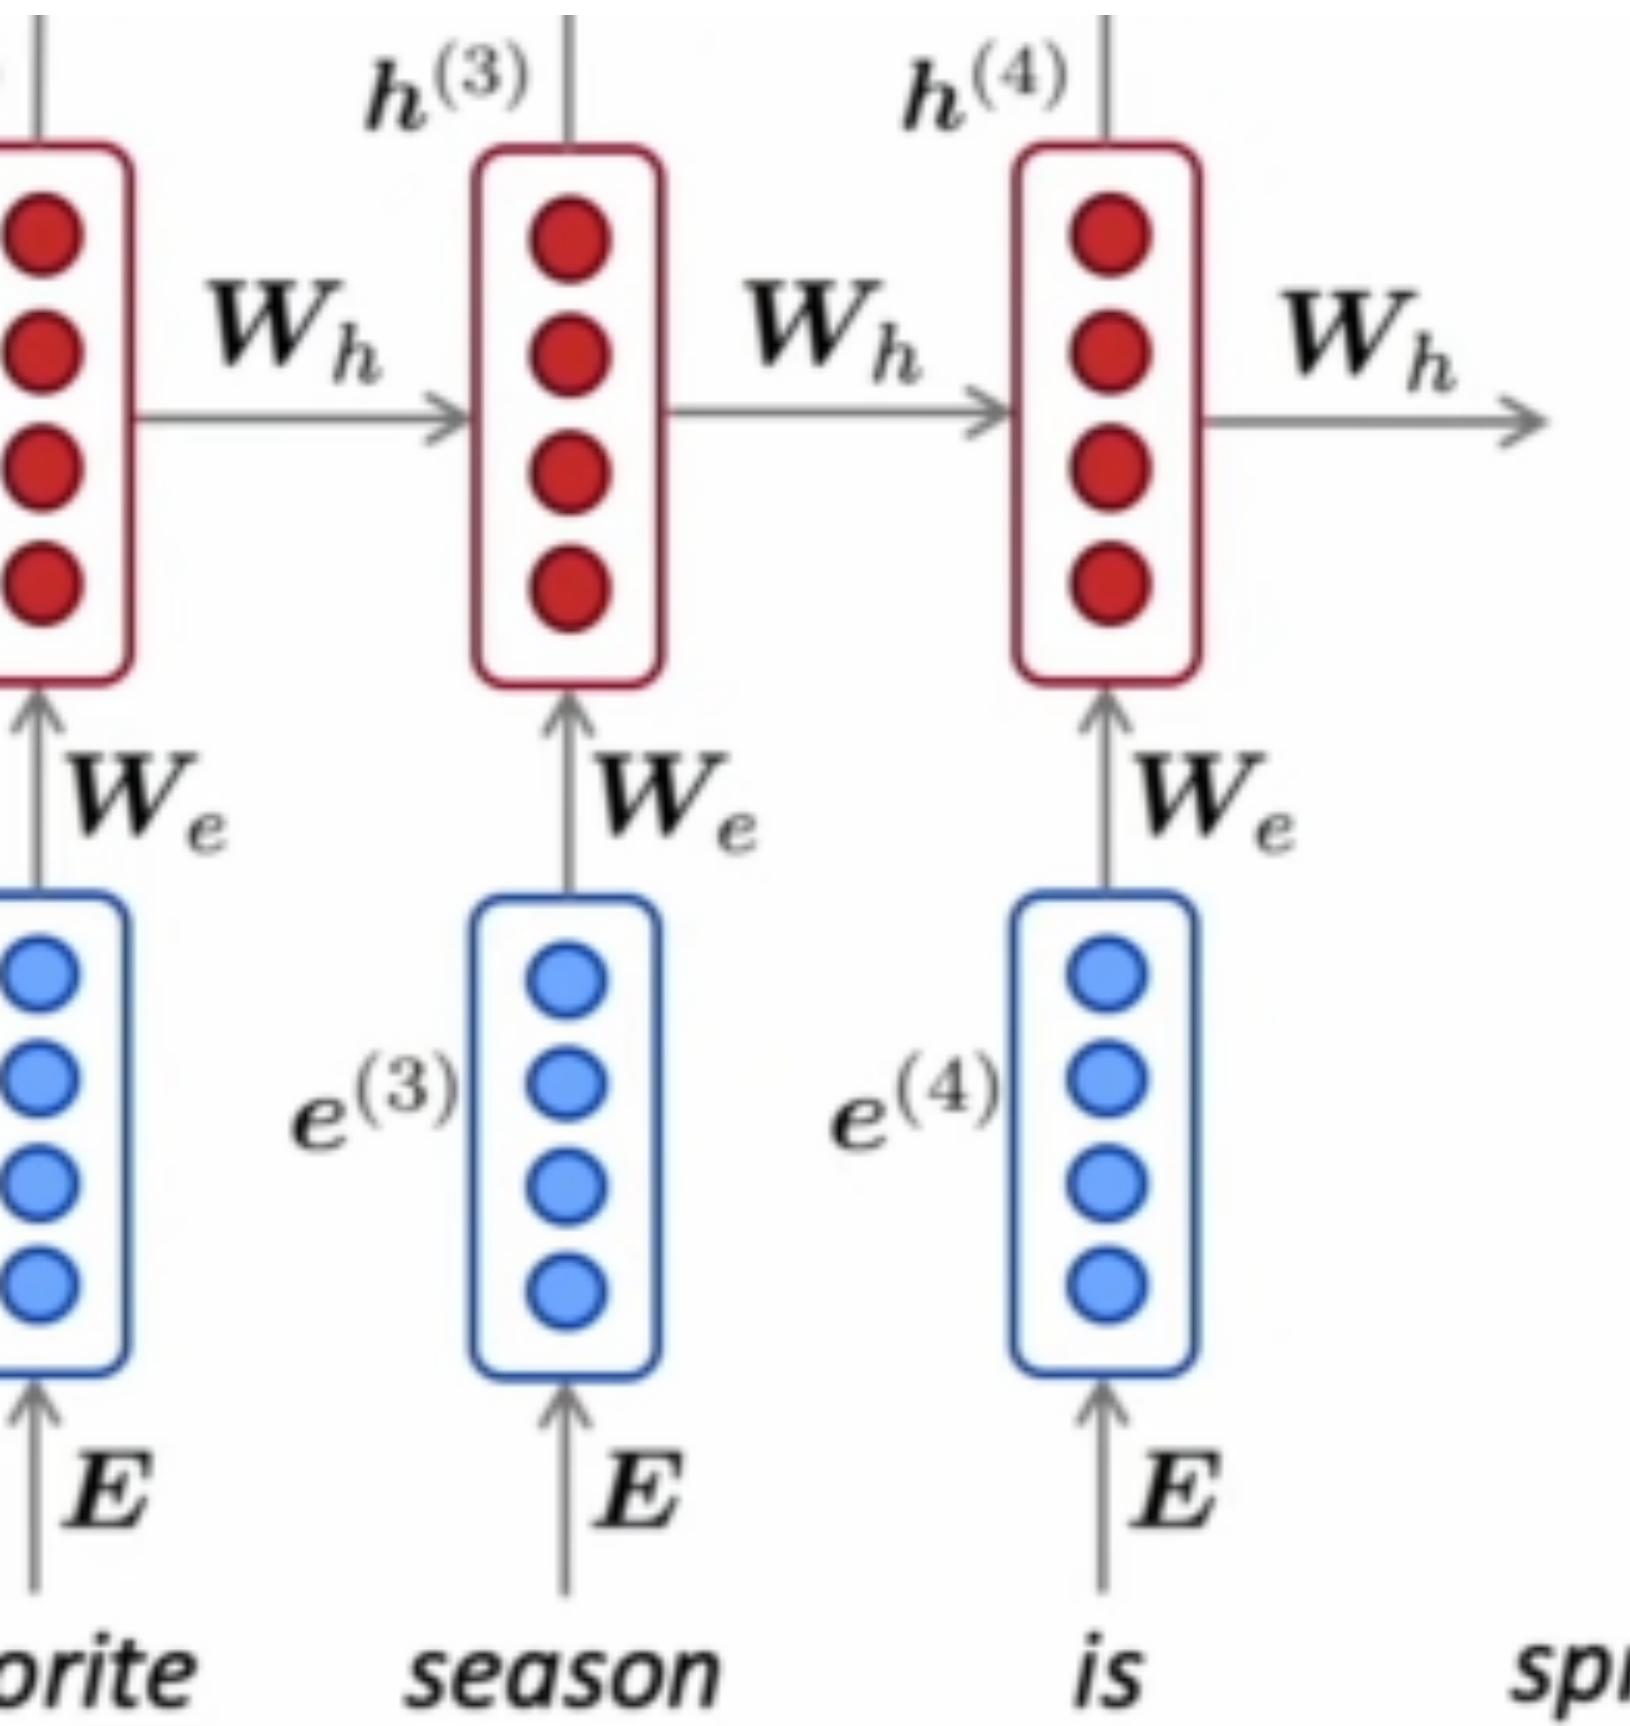
Language Model to  
next step's input.



36

my

favo



\*\*\*

*ring*

# Generating text

- Let's have some fun!
- You can train a RNN in that style.
- RNN-LM trained on Harry Potter

“Sorry,” Harry shouted,

# with a RNN Language Model!

n!

N-LM on any kind of text, the

in *Harry Potter*:



, panicking—"I'll leave those broom

# ge Model

hen generate text



ns in London, are

they?"

"No idea," said Nearly Headless Nick.  
last bit of treacle Chariot  
common room perched  
spider hadn't felt it see

**Source:** [http://www.jkrowling.com/text.html](#)

Headless Nick, casting low close by  
ms, from Harry's shoulder, and to a  
d upon it, four arms held a shining  
emed. He reached the teams too.

<https://medium.com/deep-writing/harry-potter-written>

**with a RNN Language**

Cedric, carrying the  
answer him the  
knob from when the

[-by-artificial-intelligence-8a9431803da6](#)

ge Model

- Let's have some fun
- You can train a RNN in that style.
- RNN-LM trained on

Title: CHOCOLATE RANCH

Categories: Game, Casseroles

Yield: 6 Servings

2 tb Parmesan cheese --

1 c Coconut milk

2 Eggs beaten

n!

N-LM on any kind of text, th

## recipes:

BARBECUE

, Cookies, Cookies

- chopped



en generate text



Place each pasta over layer until firm. Serve hot in bowls.

Combine the cheese and salt and stir in the chocolate and

rs of lumps. Shape mixture into the mo  
died fresh, mustard, orange and cheese  
together the dough in a large skillet  
nd pepper.

Source: <https://gist.github.com>

with a RNN Language

derate oven and simmer

e.

t; add the ingredients

<https://www.reddit.com/nylki/1efbaa36635956d35bcc>

**Model**

- Let's have some fun
- You can train a RNN in that style.
- RNN-LM trained on



! -LM on any kind of text, the

## paint color names:

31 137 165

 Sand Dan 201 17

51 124 112

 Grade Bat 48 94

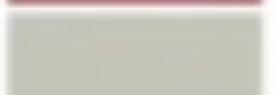
9 173 140

 Light Of Blast 17

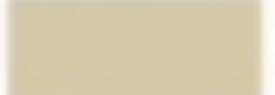
ite 221 215 236

 Grass Bat 176 99

78 181 196

 Sindis Poop 204

wn 133 104 85

 Dope 219 209 17

n 144 106 74

 Testing 156 101

n 237 217 177

 Stoner Blue 152

232 223 215

 Bubble Simp 226

n generate text

72 143

83

75 150 147

9 108

205 194

79

106

165 159

5 181 132

Burf Pink 223

Rose Hork 230

This is an example of a **char**

40

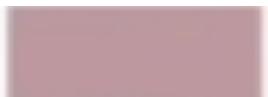
Source: [\[link\]](#)

**Why should we do it?**

- Language Modeling

173 179

0 215 198



Stanky Bean 197



Turdly 190 164

acter-level RNN-LM (predicts what ch

<http://aiweirdness.com/post/160776374467/new-paint>

care about Language

is a benchmark task that he

7 162 171

116

character comes next)

[6-colors-invented-by-neural-network](#)

e Modeling?

helps us